

What 5,000 Hours of LLM Practice Actually Looks Like

Noah T. G. Chelednik | noahchelednik@proton.me | github.com/noah-chelednik

The Short Version

I have been using ChatGPT and Claude as primary working tools since August 2023. Over 33 months I have exchanged 23.4 million tokens across 2,398 conversations on 883 active days. OpenAI confirmed I am in the top 1% of ChatGPT users by messages sent. I built an open source analytics pipeline to prove the hours claim from raw export data, published the methodology, and open sourced the scripts.

What I Build With It

The biggest project is a production domain-specific AI pipeline. 100,000+ lines of Python. 500+ modules. Hundreds of processed works across multiple source domains. Millions of structured output records. A controlled output specification with formal rules. 44 automated quality layers. Multi-format output generation. AI training data bundles with regulatory compliance documentation. IP defensibility infrastructure.

Built in four months, alone, on a single 4090. I did not write a single line of the code by hand. Every line was generated through directed collaboration with Claude.

Other projects include VCAT (Voynich Manuscript structured data, GitHub and HuggingFace, 100% test coverage), The Muser (AI music orchestration), and the LLM Export Analytics pipeline itself.

Why It Works

My data shows something that contradicts standard prompt engineering advice. 57% of my messages to the AI are informational. Not commands, not questions. Context: the problem space, the constraints, what exists, what I have tried. Only 3.3% of my messages are directive.

Most people tell the AI what to do. I tell it what the situation is and let it work within that space. The result is 6.3 words of usable output for every word I invest, with a 7.2% correction rate. Dense context constrains the model's generation space. The model does not need to guess. The constraint does the work.

I traced this pattern to my training. History, Political Science, a year of law school, Mock Trial. All of those teach the same thing: construct the complete case, present it to the decision maker, let the reasoning happen. That turns out to be the right paradigm for a system whose reasoning you cannot directly control.

The Numbers

Total tokens exchanged	23.4 million
Conversations	2,398 across ChatGPT and Claude
Active days (of ~1,000 possible)	883 (88%)
Documented in chat hours (auditable)	3,737
Total practice hours (with offline)	6,465 (conservative floor 5,101)
Models used	21 (GPT 3.5 through GPT 5.4 Thinking, o1, o3)
Total subscription cost, 33 months	\$2,080
Productive output per dollar	6,241 words

All figures reproducible from complete platform exports using open source scripts at github.com/noah-chelednik/llm-export-analytics.